Abstract

Neural machine translation (NMT) has achieved remarkable progress in recent years. However, translating long documents remains challenging because models must maintain consistency across sentences and paragraphs.

In this paper we study document-level translation with a decoder-only architecture, and evaluate three strategies for injecting context: sentence-level batching, sliding windows, and retrieval of similar segments.

Experiments

We evaluate on three language pairs: English-Chinese, English-Japanese and Chinese-Japanese. All models are quantized to four bits and served locally.

Results indicate that a sliding window of two preceding sentences improves terminology consistency by 6.8 percent without measurable latency cost.